

EXECUTIVE OVERVIEW

The DevOps OverSight Agent

An AI agent that diagnoses production incidents in minutes — across both your logging and monitoring platforms — **and never acts without a human's sign-off.**

Solution 1 · Ballerina + MCP Proxy

Solution 2 · LangChain + A2A

Measured on a production-grade AI model



Executive summary

- **The problem:** engineers lose 20–40 minutes per major incident manually cross-referencing two separate tools before they even know what's wrong.
- **What we built:** an AI agent that does that cross-tool detective work automatically — then proposes one specific fix and waits for a human to approve it.
- **It never acts alone:** no fix runs, no system changes, without an explicit human approval. This is enforced, not just policy.
- **It's measured, not promised:** on a production-grade AI model, it correctly diagnosed the test incident 18 times out of 18.

100%

correct diagnosis rate (measured)

14–36s

time to a proposed fix

0

vendor lock-in

◆ Bottom line

This isn't a concept — it's a working system with measured, repeatable results.

The cost of manually correlating two tools

- **Two systems, no shared language.** Your logs live in one platform, your metrics and traces in another. They don't talk to each other.
- **A person is the integration.** During an incident, an engineer manually copies identifiers between tools, cross-referencing by hand and by memory.
- **It doesn't scale.** The bigger and more complex your systems get, the worse this gets — and it depends on whoever happens to be on call.
- **The stakes are real.** Every minute of manual correlation is a minute of extended downtime, customer impact, and risk exposure.

20–40 min


**lost per major incident
to manual correlation alone**

◆ **Why it matters at your level**


This is a direct, measurable line item in mean-time-to-resolution — and MTTR is a direct line item in customer trust and operational cost.

What it is, in plain terms

- **It watches.** The agent continuously has access to both your logging and monitoring platforms — the same data your engineers already look at.
- **It investigates automatically.** When something breaks, it pulls evidence from both tools and connects them using the identifier that already links them together.
- **It proposes — it doesn't act.** It names one specific fix from a pre-approved list. It never invents an action and never runs one unsupervised.
- **A human always decides.** The fix only executes after an engineer explicitly approves it. Every action is logged.



The workflow cycle



01

TWO WAYS WE BUILT IT

We built it twice, to compare

Same capability, same test scenario — two different engineering approaches, so we could evaluate real trade-offs instead of guessing.

SOLUTION 1

Ballerina + MCP Proxy

One agent, one gateway to both platforms. Built to run natively on our WSO2 Agent Manager platform.


SOLUTION 2

LangChain + A2A

Specialist agents divide the work, one per platform. Built in Python — the most widely adopted AI agent ecosystem.


Ballerina + MCP Proxy

- **One agent, one gateway.** A single AI agent connects through one gateway that reaches both your logging and monitoring platforms.
- **Fewer moving parts.** Fewer components to operate, secure, and audit — a simpler system to run day to day.
- **Native to our platform.** Runs directly on WSO2's own Agent Manager, our platform for governing and observing AI agents.
- **One consistent language.** The entire stack — agent, gateway, and services — is built in a single language end to end.



LangChain + A2A (Agent to Agent)

- **Specialist agents.** A lead agent delegates to two specialist agents, one dedicated to each platform.
- **Built on Python.** Uses LangChain, the most widely adopted framework for building AI agents — a large hiring and support pool.
- **Approval enforced in code.** The human-approval step is built into the system's control flow, not just an instruction to the AI.
- **Scales to many platforms.** This division-of-labor pattern extends cleanly if more tools or teams are added later.



What we actually measured

Metric — measured on a production-grade AI model	Ballerina + Proxy	LangChain + A2A
Correct diagnosis rate	100% (18 / 18 test runs)	100% (18 / 18 test runs)
Median time to a proposed fix	~14 seconds	~36 seconds
Test runs evaluated	18 (3 independent batches of 6)	18 (3 independent batches of 6)

◆ Why this matters

Both are reliable enough to trust in a live environment. Ballerina resolves faster in this test because it makes fewer hops between agents for the same investigation.

◆ A critical cost finding

A free, locally-hosted AI model is possible for either approach, but was only 28–61% reliable in the same test. Reliability came from using a production-grade model, not from which architecture we chose — worth weighing directly against licensing cost.

Method: identical simulated incident injected into a test service; timed end-to-end from alert to a proposed fix; 3 independent test batches of 6 runs per approach. Testing cost using Claude Haiku 4.5: \$2.56.

Nothing changes without a human

1

Propose, never act alone

The agent names a fix and stops. It cannot execute a change without an explicit, recorded human approval.

2

A bounded set of actions

The agent can only ever propose from a fixed, pre-approved list of fixes — never an arbitrary command.

3

Every action is logged

Every proposal, approval, and executed action is recorded in an audit trail.

No AI vendor lock-in

- **The AI model is a setting, not a commitment.** Switch between AI providers, or run entirely on your own infrastructure, without changing the system.
- **Works fully offline if required.** A locally-hosted, no-cost model option exists for environments that can't use an outside AI vendor.
- **Same investigation, any provider.** We've proven the same investigation completes successfully across multiple AI providers with no code changes.

Local model

no cost, offline

Cloud providers

your choice of vendor

WSO2 gateway

governed, audited access

A real 5-minute demonstration

1

Inject

A real fault is introduced into a test service.

2

Detect & investigate

The agent notices and automatically investigates across both platforms.

3

Propose & approve

It proposes a specific fix; an operator approves it.

4

Recover

The fix runs and the system returns to healthy — automatically documented.

◆ Runs on a laptop, no live accounts required

This can be demonstrated live, end to end, in under five minutes — including for a client, on the spot.

Our read, and what would decide it

Both approaches are production-viable on a real AI model.
Ballerina measured faster; LangChain offers a code-enforced approval gate and a larger talent pool.

Favors Ballerina + Proxy

- Standardizing on our own WSO2 platform.
- Fewest components to run and secure.
- Faster measured response time.

Favors LangChain + A2A

- Teams already build in Python.
- Approval gate enforced in code, by design.
- Broadest hiring and support pool.

Path to a production pilot

1 Choose a direction — or run a short pilot of both against a real, low-risk service.

2 Connect it to your live logging and monitoring platforms (no code changes required).

3 Finalize the approval workflow and audit requirements with your compliance team.

4 Run a live, supervised pilot on a non-critical service before wider rollout.

02




SECTION


Reference Deployment Architecture

Putting It All Together


WSO2 Agent Platform: Agent Manager



OpenChoreo on SUSE Rancher



Agent Platform + DevOps Oversight Agent



Thank you.

Happy to go deeper on any part of this — architecture, results, or rollout planning.

Scott Bechtel · WSO2 Forward Deployed Engineering

